

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
```

```
a=pd.read_csv("/content/WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

```
a
```



	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educa
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2
...	...	...	...	...	...		...
1465	36	No	Travel_Frequently	884	Research & Development		23
1466	39	No	Travel_Rarely	613	Research & Development		6
1467	27	No	Travel_Rarely	155	Research & Development		4
1468	49	No	Travel_Frequently	1023	Sales		2
1469	34	No	Travel_Rarely	628	Research & Development		8

1470 rows × 35 columns

## Read the data types

```
a.dtypes
```

```
Age                int64
Attrition          object
BusinessTravel     object
DailyRate         int64
Department        object
DistanceFromHome  int64
Education         int64
EducationField     object
EmployeeCount     int64
EmployeeNumber    int64
EnvironmentSatisfaction int64
Gender            object
HourlyRate        int64
JobInvolvement    int64
JobLevel          int64
JobRole           object
JobSatisfaction   int64
MaritalStatus     object
MonthlyIncome     int64
MonthlyRate       int64
NumCompaniesWorked int64
Over18            object
OverTime          object
PercentSalaryHike int64
PerformanceRating int64
RelationshipSatisfaction int64
StandardHours     int64
StockOptionLevel  int64
TotalWorkingYears int64
TrainingTimesLastYear int64
WorkLifeBalance   int64
YearsAtCompany    int64
YearsInCurrentRole int64
YearsSinceLastPromotion int64
YearsWithCurrManager int64
dtype: object
```

```
a.shape
```

(1470, 35)

```
a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1470 non-null   int64
1   Attrition                            1470 non-null   object
2   BusinessTravel                        1470 non-null   object
3   DailyRate                            1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                             1470 non-null   int64
7   EducationField                        1470 non-null   object
8   EmployeeCount                         1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                       1470 non-null   int64
17  MaritalStatus                        1470 non-null   object
18  MonthlyIncome                        1470 non-null   int64
19  MonthlyRate                          1470 non-null   int64
20  NumCompaniesWorked                   1470 non-null   int64
21  Over18                               1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                    1470 non-null   int64
24  PerformanceRating                    1470 non-null   int64
25  RelationshipSatisfaction              1470 non-null   int64
26  StandardHours                        1470 non-null   int64
27  StockOptionLevel                     1470 non-null   int64
28  TotalWorkingYears                    1470 non-null   int64
29  TrainingTimesLastYear                1470 non-null   int64
30  WorkLifeBalance                      1470 non-null   int64
31  YearsAtCompany                       1470 non-null   int64
32  YearsInCurrentRole                   1470 non-null   int64
33  YearsSinceLastPromotion               1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

Statistics about the dataset

```
a.describe()
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	Employee
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024
std	9.135373	403.509100	8.106864	1.024165	0.0	602
min	18.000000	102.000000	1.000000	1.000000	1.0	1
25%	30.000000	465.000000	2.000000	2.000000	1.0	491
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555
max	60.000000	1499.000000	29.000000	5.000000	1.0	2065

8 rows × 6 columns

Null values identification

```
a.isnull().any()
```

Age	False
Attrition	False
BusinessTravel	False
DailyRate	False
Department	False
DistanceFromHome	False

```

Education                False
EducationField            False
EmployeeCount             False
EmployeeNumber            False
EnvironmentSatisfaction   False
Gender                    False
HourlyRate                False
JobInvolvement            False
JobLevel                  False
JobRole                   False
JobSatisfaction           False
MaritalStatus             False
MonthlyIncome             False
MonthlyRate               False
NumCompaniesWorked        False
Over18                    False
OverTime                  False
PercentSalaryHike         False
PerformanceRating         False
RelationshipSatisfaction   False
StandardHours             False
StockOptionLevel          False
TotalWorkingYears         False
TrainingTimesLastYear     False
WorkLifeBalance           False
YearsAtCompany            False
YearsInCurrentRole        False
YearsSinceLastPromotion   False
YearsWithCurrManager      False
dtype: bool

```

```
a.isnull().sum()
```

```

Age                      0
Attrition                0
BusinessTravel           0
DailyRate               0
Department              0
DistanceFromHome         0
Education                0
EducationField            0
EmployeeCount            0
EmployeeNumber           0
EnvironmentSatisfaction   0
Gender                   0
HourlyRate               0
JobInvolvement           0
JobLevel                 0
JobRole                  0
JobSatisfaction           0
MaritalStatus            0
MonthlyIncome            0
MonthlyRate              0
NumCompaniesWorked        0
Over18                   0
OverTime                 0
PercentSalaryHike         0
PerformanceRating         0
RelationshipSatisfaction   0
StandardHours             0
StockOptionLevel          0
TotalWorkingYears         0
TrainingTimesLastYear     0
WorkLifeBalance           0
YearsAtCompany            0
YearsInCurrentRole        0
YearsSinceLastPromotion   0
YearsWithCurrManager      0
dtype: int64

```

## Data Visualization

```

d=a.corr()
d

```

```
<ipython-input-13-385900cf86c7>:1: FutureWarning: The default value of numeric_only i
d=a.corr()
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCo
<b>Age</b>	1.000000	0.010661	-0.001686	0.208034	↑
<b>DailyRate</b>	0.010661	1.000000	-0.004985	-0.016806	↑
<b>DistanceFromHome</b>	-0.001686	-0.004985	1.000000	0.021042	↑
<b>Education</b>	0.208034	-0.016806	0.021042	1.000000	↑
<b>EmployeeCount</b>	NaN	NaN	NaN	NaN	↑
<b>EmployeeNumber</b>	-0.010145	-0.050990	0.032916	0.042070	↑
<b>EnvironmentSatisfaction</b>	0.010146	0.018355	-0.016075	-0.027128	↑
<b>HourlyRate</b>	0.024287	0.023381	0.031131	0.016775	↑
<b>JobInvolvement</b>	0.029820	0.046135	0.008783	0.042438	↑
<b>JobLevel</b>	0.509604	0.002966	0.005303	0.101589	↑
<b>JobSatisfaction</b>	-0.004892	0.030571	-0.003669	-0.011296	↑
<b>MonthlyIncome</b>	0.497855	0.007707	-0.017014	0.094961	↑
<b>MonthlyRate</b>	0.028051	-0.032182	0.027473	-0.026084	↑
<b>NumCompaniesWorked</b>	0.299635	0.038153	-0.029251	0.126317	↑
<b>PercentSalaryHike</b>	0.003634	0.022704	0.040235	-0.011111	↑
<b>PerformanceRating</b>	0.001904	0.000473	0.027110	-0.024539	↑
<b>RelationshipSatisfaction</b>	0.053535	0.007846	0.006557	-0.009118	↑
<b>StandardHours</b>	NaN	NaN	NaN	NaN	↑
<b>StockOptionLevel</b>	0.037510	0.042143	0.044872	0.018422	↑
<b>TotalWorkingYears</b>	0.680381	0.014515	0.004628	0.148280	↑
<b>TrainingTimesLastYear</b>	-0.019621	0.002453	-0.036942	-0.025100	↑

```
plt.subplots(figsize=(15,15))
sns.heatmap(d,annot=True)
```

&lt;Axes: &gt;



```
f = plt.figure()
f.set_figwidth(15)
f.set_figheight(12)
```

```
# Subplot 1
plt.subplot(3, 3, 1)
sns.countplot(x="Attrition", data=a)
```

```
# Subplot 2
plt.subplot(3, 3, 2)
sns.countplot(x="BusinessTravel", data=a)
```



```
# Subplot 5
plt.subplot(3, 3, 3)
sns.countplot(x="Department", data=a)
```

```
# Subplot 8
plt.subplot(3, 3, 4)
sns.countplot(x="EducationField", data=a)
```

```
# Subplot 9
plt.subplot(3, 3, 5)
sns.countplot(x="OverTime", data=a)
```

```
# Adjust layout
plt.tight_layout()
```


```
# Show the plots
plt.show()
```



Outlier Detection


```
sns.boxplot(x="Age",data=a)
```

<Axes: xlabel='Age'>




```
sns.boxplot(x="DailyRate",data=a)
```

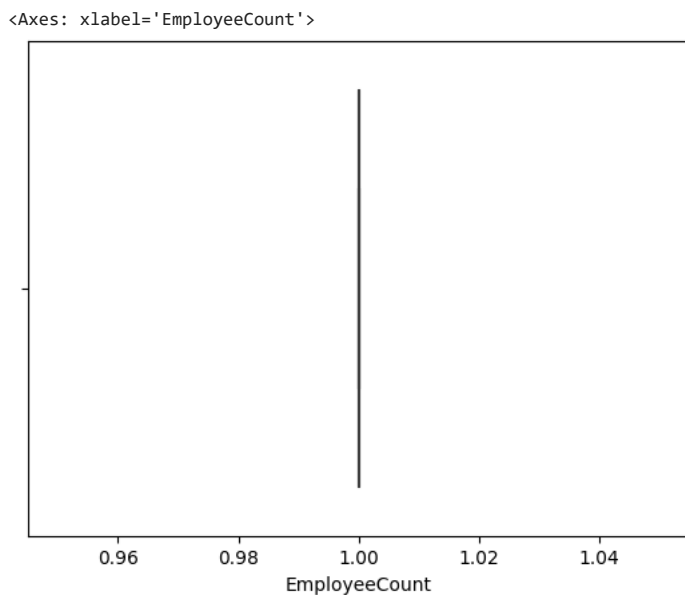
```
<Axes: xlabel='DistanceFromHome'>  
sns.boxplot(x="DistanceFromHome", data=a)
```



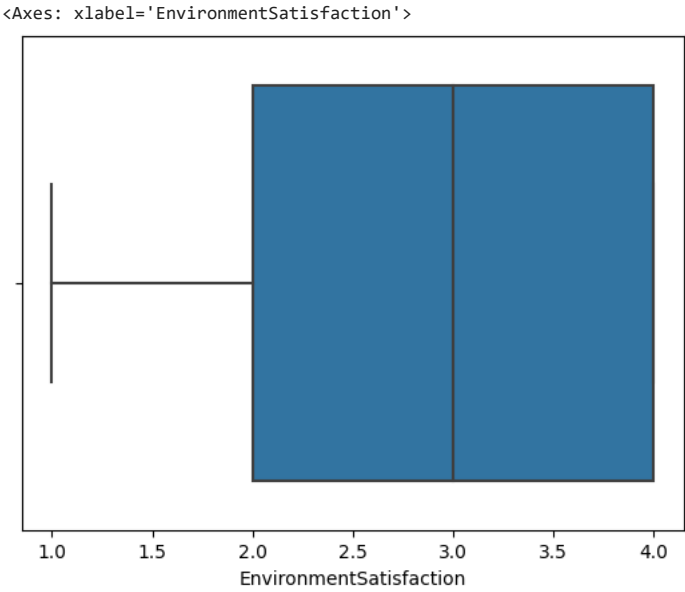
```
sns.boxplot(x="Education", data=a)
```



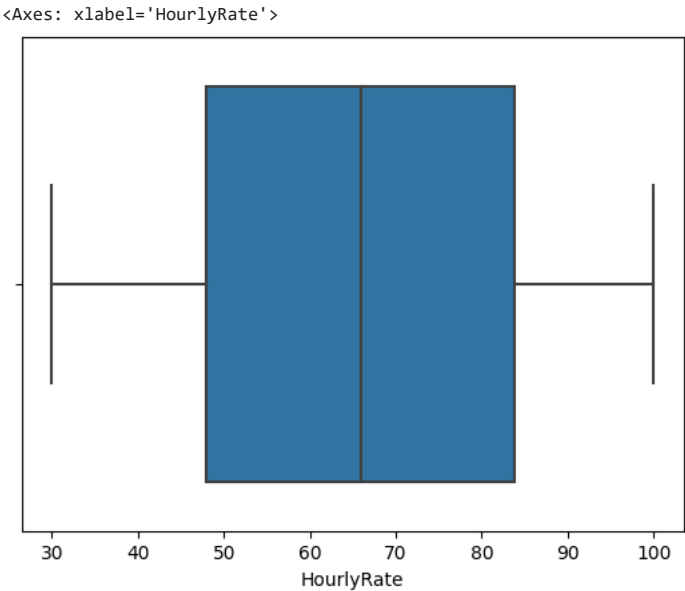
```
sns.boxplot(x="EmployeeCount", data=a)
```



```
sns.boxplot(x="EnvironmentSatisfaction",data=a)
```



```
sns.boxplot(x="HourlyRate",data=a)
```



Splitting dependent and independent variables

```
x=a.drop(columns=["Attrition"],axis=1)
x.head()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction
0	41	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	3.0
1	49	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	3.0
2	37	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	3.0
3	33	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	3.0
4	27	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	3.0

5 rows × 34 columns

```
x.shape
```

(1470, 34)



```
y=a["Attrition"]
y.head()

0    Yes
1    No
2    Yes
3    No
4    No
Name: Attrition, dtype: object
```

```
y.shape

(1470,)
```

Encoding

```
from sklearn.preprocessing import LabelEncoder
```

```
l=LabelEncoder()
```

```
x["Gender"]=l.fit_transform(x["Gender"])
x['Gender']
```

```
0      0
1      1
2      1
3      0
4      1
..
1465    1
1466    1
1467    1
1468    1
1469    1
Name: Gender, Length: 1470, dtype: int64
```

```
x['Gender'].value_counts()
```

```
1    882
0    588
Name: Gender, dtype: int64
```

```
x['Gender'].nunique()
```

```
2
```

```
x.head()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Environment
0	41	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	
1	49	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	
2	37	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	
3	33	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	
4	27	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	

5 rows × 11 columns

```
Dept = pd.get_dummies(a, columns=["Department"])
print(Dept)
```

```

4          3          4          4
...      ...      ...      ...
1465      3          5          2
1466      3          7          7
1467      3          6          2
1468      2          9          6
1469      4          4          3

```

```

YearsSinceLastPromotion YearsWithCurrManager \
0          0          5
1          1          7
2          0          0
3          3          0
4          2          2
...      ...      ...
1465      0          3
1466      1          7
1467      0          3
1468      0          8
1469      1          2

```

```

Department_Human Resources Department_Research & Development \
0          0          0
1          0          1
2          0          1
3          0          1
4          0          1
...      ...      ...
1465      0          1
1466      0          1
1467      0          1
1468      0          0
1469      0          1

```

```

Department_Sales
0          1
1          0
2          0
3          0
4          0
...      ...
1465      0
1466      0
1467      0
1468      1
1469      0

```

[1470 rows x 37 columns]

print(x)

```

1466      6          1      Medical      1
1467      4          3      Life Sciences 1
1468      2          3      Medical      1
1469      8          3      Medical      1

```

```

EmployeeNumber EnvironmentSatisfaction ... RelationshipSatisfaction \
0          1          2 ...          1
1          2          3 ...          4
2          4          4 ...          2
3          5          4 ...          3
4          7          1 ...          4
...      ...      ...      ...
1465      2061      3 ...          3
1466      2062      4 ...          1
1467      2064      2 ...          2
1468      2065      4 ...          4
1469      2068      2 ...          1

```

1466	5	3	/
1467	0	3	6
1468	3	2	9
1469	3	4	4

	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2
...	...	...	...
1465	2	0	3
1466	7	1	7
1467	2	0	3
1468	6	0	8
1469	3	1	2

[1470 rows x 34 columns]

a.head()

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educatic
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2

5 rows x 35 columns

x.head()

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Environm
0	41	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	
1	49	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	
2	37	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	
3	33	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	
4	27	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	

5 rows x 34 columns

```
Dept=pd.get_dummies(x["Department"],drop_first=True)
Dept
```

```
x=pd.concat([x,Dept],axis=1)

x.head()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Environm
0	41	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	
1	49	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	
2	37	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	
3	33	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	
4	27	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	

5 rows × 38 columns

Feature Scaling

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X = a[['Age', 'MonthlyIncome', 'YearsAtCompany', 'JobSatisfaction', 'EnvironmentSatisfaction', 'YearsWithCurrManager', 'WorkLifeBalance']]
Y = a['Attrition']

X.head()
```

	Age	MonthlyIncome	YearsAtCompany	JobSatisfaction	EnvironmentSatisfaction	YearsWithCurrManager	WorkLifeBalance
0	41	5993	6	4	2	5	1
1	49	5130	10	2	3	7	3
2	37	2090	0	3	4	0	3
3	33	2909	8	3	4	0	3
4	27	3468	2	2	1	2	3

```
x.tail()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Envir
1465	36	Travel_Frequently	884	Research & Development	23	2	Medical	1	2061	
1466	39	Travel_Rarely	613	Research & Development	6	1	Medical	1	2062	
1467	27	Travel_Rarely	155	Research & Development	4	3	Life Sciences	1	2064	
1468	49	Travel_Frequently	1023	Sales	2	3	Medical	1	2065	
1469	34	Travel_Rarely	628	Research & Development	8	3	Medical	1	2068	

5 rows × 38 columns

x

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Envir
0	41	Travel_Rarely	1102	Sales		1	2	Life Sciences	1	1
1	49	Travel_Frequently	279	Research & Development		8	1	Life Sciences	1	2
2	37	Travel_Rarely	1373	Research & Development		2	2	Other	1	4
3	33	Travel_Frequently	1392	Research & Development		3	4	Life Sciences	1	5
4	27	Travel_Rarely	591	Research & Development		2	1	Medical	1	7
...	...	...	...	...		...	...	...	...	...
1465	36	Travel_Frequently	884	Research & Development		23	2	Medical	1	2061
1466	39	Travel_Rarely	613	Research & Development		6	1	Medical	1	2062

Splitting data into test and train

1468	49	Travel_Frequently	1023	Sales		2	3	Medical	1	2065
------	----	-------------------	------	-------	--	---	---	---------	---	------

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

X\_train,X\_test,Y\_train,Y\_test.shape

860	22	2853	0	4
1459	29	4025	4	2
1126	50	19331	1	3
EnvironmentSatisfaction    YearsWithCurrManager    WorkLifeBalance				
1097		3	0	3
727		2	0	3
254		4	2	3
1175		4	0	3
1341		2	2	3
...	...	...	...	...
1130		2	8	2
1294		2	2	1
860		3	0	3
1459		4	3	3
1126		3	0	3
[1176 rows x 7 columns],				
Age    MonthlyIncome    YearsAtCompany    JobSatisfaction \				
1041	28	8463	5	1
184	53	4450	4	1
1222	24	1555	1	3
67	45	9724	1	1
220	36	5914	13	2
...	...	...	...	...
567	34	6274	6	4
560	34	5121	0	1
945	50	16880	3	1
522	37	4680	1	4
651	47	4537	7	4
EnvironmentSatisfaction    YearsWithCurrManager    WorkLifeBalance				
1041		4	3	3
184		4	3	3
1222		4	0	3
67		2	0	3
220		4	7	4
...	...	...	...	...
567		4	4	3
560		2	0	3
945		4	2	3
522		4	0	3
651		3	7	3
[294 rows x 7 columns],				
1097	No			
727	No			

```
Name: Attrition, Length: 1176, dtype: object,
(294,))
```

## Logistic Regression

## Model Building & Import the model building Libraries

```
from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
```

```
model.fit(X_train, Y_train)
```

- ▼ LogisticRegression

```
LogisticRegression()
```

```
pred=model.predict(X_test)
```

pred

[illegible]

Y\_test

```
1041    No
184     No
1222    Yes
67      No
220     No
...
567     No
560     No
945     No
522     No
651     No
Name: Attrition, Length: 294, dtype: object
```

a

	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber
1	Yes	Travel_Rarely	1102	Sales		1	2	Life Sciences	1
2	No	Travel_Frequently	279	Research & Development		8	1	Life Sciences	1
3	Yes	Travel_Rarely	1373	Research & Development		2	2	Other	1
4	No	Travel_Frequently	1392	Research & Development		3	4	Life Sciences	1
5	No	Travel_Rarely	591	Research & Development		2	1	Medical	1

Evaluation of classification model

```
#Accuracy score
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report,roc_auc_score,roc_curve
```

```
accuracy = accuracy_score(Y_test, pred)
```

```
report = classification_report(Y_test, pred, zero_division=1)
```

```
print(f'Accuracy: {accuracy}')
print(f'Classification Report:\n{report}')
```

Accuracy: 0.8673469387755102				
Classification Report:				
	precision	recall	f1-score	support
No	0.87	1.00	0.93	255
Yes	1.00	0.00	0.00	39
accuracy			0.87	294
macro avg	0.93	0.50	0.46	294
weighted avg	0.88	0.87	0.81	294

```
confusion_matrix(Y_test,pred)
```

array([[255, 0],
[ 39, 0]])

```
pd.crosstab(Y_test,pred)
```

col_0	No
Attrition	
No	255
Yes	39

Roc-AUC curve

```
probability=model.predict_proba(X_test)[: ,1]
```

```
probability
```

array([0.14873939, 0.17373604, 0.25084589, 0.1865791 , 0.11911736,
0.14963007, 0.15969356, 0.20644099, 0.08193936, 0.18537088,
0.16096129, 0.02189805, 0.15660552, 0.11782876, 0.18248771,
0.13287268, 0.14334387, 0.0892007 , 0.06858367, 0.05708061,
0.1753651 , 0.14395111, 0.10012064, 0.15057687, 0.2329628 ,
0.03338823, 0.27116899, 0.15771847, 0.18762417, 0.10029771,
0.10548668, 0.15048832, 0.12644386, 0.14778903, 0.2030313 ,
0.06737083, 0.04935137, 0.35253675, 0.19926437, 0.23846212,
0.08198467, 0.28864726, 0.23955634, 0.19282515, 0.22246873,
0.11288909, 0.17545014, 0.24051176, 0.14059822, 0.32377579,
0.08977525, 0.15148043, 0.01896052, 0.14635136, 0.20158982,
0.10191406, 0.10573264, 0.08537077, 0.1631479 , 0.12443613,
0.10510977, 0.33623452, 0.11027653, 0.05493965, 0.28005007,
0.18450873, 0.12499531, 0.17197795, 0.17873294, 0.06110176,
0.18127058, 0.08791989, 0.15005295, 0.15959692, 0.19866202,
0.07388538, 0.19341696, 0.19100387, 0.08712656, 0.08033949,
0.02928375, 0.13253218, 0.05956382, 0.16844953, 0.08753921,
0.17957672, 0.12899389, 0.16872069, 0.16947305, 0.12397644,
0.1099147 , 0.24576674, 0.07821105, 0.2716565 , 0.12140547,

```

0.06524951, 0.1337184 , 0.14536957, 0.18726004, 0.10915274,
0.04570312, 0.10169758, 0.07390408, 0.22704117, 0.07208355,
0.08035364, 0.18593691, 0.16647288, 0.10818369, 0.05315879,
0.17696614, 0.18973955, 0.22476227, 0.17342537, 0.21403334,
0.16943373, 0.16771766, 0.09747364, 0.11387728, 0.2559594 ,
0.32393512, 0.08431327, 0.13118746, 0.10751731, 0.09837008,
0.25991497, 0.18954525, 0.11954205, 0.10534474, 0.09694665,
0.07268098, 0.30507638, 0.06501248, 0.14080365, 0.1255734 ,
0.11537899, 0.23299235, 0.17264787, 0.24765337, 0.06927027,
0.21512755, 0.09901074, 0.16646941, 0.08047622, 0.03233445,
0.15363939, 0.14131117, 0.25851265, 0.26761484, 0.1665985 ,
0.10685997, 0.11549038, 0.19827264, 0.19076354, 0.13247131,
0.26173972, 0.17180386, 0.21324175, 0.04115976, 0.15054569,
0.16012435, 0.09434315, 0.09921354, 0.22000675, 0.06421677,
0.16643204, 0.12016002, 0.14827189, 0.08450615, 0.05725373,
0.12102272, 0.02681568, 0.18300015, 0.21076054, 0.11715199,
0.16127828, 0.18483891, 0.09043029, 0.14086669, 0.20253644,
0.0594472 , 0.10383826, 0.01617733, 0.15428555, 0.08595314,
0.22434066, 0.11577713, 0.07998958, 0.07811109, 0.12006351,
0.12845942, 0.14824842, 0.10405812, 0.19816497, 0.1162661 ,
0.21477996, 0.24395257, 0.04972863, 0.2156586 , 0.16831872,
0.17867722, 0.15398516, 0.21871738, 0.03416769, 0.07072713,
0.22242289, 0.10244091, 0.10919764, 0.12517809, 0.0706504 ,
0.07399615, 0.24438034, 0.17159597, 0.17617076, 0.10663942,
0.13898632, 0.15178097, 0.10545546, 0.2723432 , 0.07462743,
0.23465253, 0.26405405, 0.10124306, 0.3028089 , 0.12410107,
0.1909214 , 0.20302625, 0.13276688, 0.0401135 , 0.18943046,
0.23129363, 0.25951761, 0.08630086, 0.21347439, 0.20469075,
0.13330949, 0.08581729, 0.10996842, 0.06690194, 0.04616928,
0.18853288, 0.11542819, 0.21231547, 0.03597583, 0.07176025,
0.17130681, 0.11593175, 0.23407496, 0.1533375 , 0.09696206,
0.16256038, 0.06366454, 0.04689748, 0.0855508 , 0.23703024,
0.07106702, 0.18067446, 0.2069784 , 0.22648723, 0.02715875,
0.17170263, 0.14167865, 0.276632 , 0.10463943, 0.12037205,
0.21133882, 0.02933273, 0.0973697 , 0.23466029, 0.23184945,
0.1882965 , 0.04906958, 0.19036583, 0.1399965 , 0.11412922,
0.22223015, 0.12517666, 0.24824295, 0.07113102, 0.07508479,
0.14609486, 0.15491467, 0.18318556, 0.09382192, 0.04811606,

```

```


from sklearn.preprocessing import LabelBinarizer
lb = LabelBinarizer()
Y_test_bin = lb.fit_transform(Y_test)
fpr, tpr, thresholds = roc_curve(Y_test_bin, probability)

```

```

plt.plot(fpr,tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE')
plt.show()

```



## Decision Tree

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report

```

```
dt_model = DecisionTreeClassifier(random_state=50)
```



```
dt_model.fit(X_train, Y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(random_state=50)
```

```
dt_predictions = dt_model.predict(X_test)
```

```
dt_accuracy = accuracy_score(Y_test, dt_predictions)
```

```
dt_report = classification_report(Y_test, dt_predictions)
```

```
print(f'Decision Tree Accuracy: {dt_accuracy}')
```

```
Decision Tree Accuracy: 0.7789115646258503
```

```
print(f'Decision Tree Classification Report:\n{dt_report}')
```

```
Decision Tree Classification Report:
              precision    recall  f1-score   support

     No         0.90        0.84        0.87        255
     Yes         0.28        0.41        0.33         39

 accuracy          0.78        0.78        0.78        294
 macro avg         0.59        0.62        0.60        294
 weighted avg         0.82        0.78        0.80        294
```

## Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_model = RandomForestClassifier(random_state=50)
```

```
rf_model.fit(X_train, Y_train)
```

```
RandomForestClassifier
RandomForestClassifier(random_state=50)
```

```
rf_predictions = rf_model.predict(X_test)
```

```
rf_accuracy = accuracy_score(Y_test, rf_predictions)
```

```
rf_report = classification_report(Y_test, rf_predictions)
```

```
print(f'Random Forest Accuracy: {rf_accuracy}')
```

```
Random Forest Accuracy: 0.8435374149659864
```

```
print(f'Random Forest Classification Report:\n{rf_report}')
```

```
Random Forest Classification Report:
              precision    recall  f1-score   support

     No         0.88        0.95        0.91        255
     Yes         0.33        0.18        0.23         39

 accuracy          0.84        0.84        0.84        294
 macro avg         0.61        0.56        0.57        294
 weighted avg         0.81        0.84        0.82        294
```